

Introduction to R for Data Analysis in the Health Sciences: Lecture 2

Amy Willis, Biostatistics, UW

11 October, 2019

Welcome back

- ▶ Check out syllabus for *scope*, *grading*, *office hours*, *course policies* and *resources*
- ▶ Look at Lecture 1 slides for critical introductory material
- ▶ Congratulations on fantastic work on In-Class Exercise and Homework 1!

Norms and ground rules

For all:

- Be respectful and understanding of our diverse experiences
 - Both our *programming experiences* and our *life experiences*
- If comfortable, share your questions and your understandings with the whole class...
 - ...so everyone can benefit from your learning
- Avoid distracting other people during class
 - This includes with gossiping, shopping online, watching TV and using Facebook

For Amy, Serge and Thayer:

- Teach and help inclusively
- Post common questions and their answers on Canvas
- Check in with the class regularly to confirm understanding/following
- Recognise different learning styles and adapt; offer multiple pathways to learning

Normalising struggle

If this is your first time programming in R, it is likely that you will experience frustration and challenge in this course. Know that this is normal and a part of the learning process!

You can do it!

Outline of today plan

- ▶ Lecture (data manipulation)
 - ▶ Quick review of last week
 - ▶ New material: data manipulation
- ▶ Break
- ▶ In-class exercise

Recap

Last week we talked about the package tidyverse...

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1
```

```
## v ggplot2 3.2.1      v purrr  0.3.2
```

```
## v tibble  2.1.3      v dplyr  0.8.3
```

```
## v tidyr   1.0.0      v stringr 1.4.0
```

```
## v readr   1.3.1      v forcats 0.4.0
```

```
## -- Conflicts ----- tidyverse_conflic
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()    masks stats::lag()
```

Recap

...and we talked about reading in and storing data...

```
fev_data <- "/Users/adwillis/teaching/19-509/datasets/fev.csv" %>%  
  read_csv
```

```
## Parsed with column specification:  
## cols(  
##   seqnbr = col_double(),  
##   subjid = col_double(),  
##   age = col_double(),  
##   fev = col_double(),  
##   height = col_double(),  
##   sex = col_double(),  
##   smoke = col_double()  
## )
```

Recap

... and we showed how we can filter data...

```
is_tall <- fev_data$height > 72  
fev_data[is_tall, ]
```

```
## # A tibble: 7 x 7  
##   seqnbr subjid   age   fev height  sex  smoke  
##   <dbl>  <dbl> <dbl> <dbl>  <dbl> <dbl> <dbl>  
## 1    401  18841    14  4.27   72.5     1     0  
## 2    450  32741    13  4.22    74     1     0  
## 3    464  37241    13  4.88    73     1     0  
## 4    517  49541    13  5.08    74     1     0  
## 5    550  59941    14  4.27   72.5     1     0  
## 6    632  37441    17  5.63    73     1     0  
## 7    636  44241    16  3.64   73.5     1     0
```


Recap

... and take some summary statistics

```
mean(fev_data[is_tall, ]$fev)
```

```
## [1] 4.572429
```

```
sd(fev_data[is_tall, ]$fev)
```

```
## [1] 0.6632754
```

Reminders

- ▶ The first line of any R script should be `library(tidyverse)`
- ▶ We refer to the rows and columns of a data frame using square brackets, e.g., `fev_data[is_tall,]`
- ▶ We call functions using `function_name(input)` or `input %>% function_name`

Functions and arguments

Functions can have multiple arguments (inputs)

```
some_heights <- c(162, 170, 176, NA)
sd(some_heights)
```

```
## [1] NA
```

```
sd(some_heights, na.rm = TRUE)
```

```
## [1] 7.023769
```

Functions and arguments

Arguments can be referred to by name, or given in the canonical order. Find the canonical order by asking:

```
?sd
```

Getting help

Some functions are well documented and give helpful information.
Others are not.

The internet is a fantastic resource!

Getting help



[r] tidyverse how do i filter na



[All](#)

[Videos](#)

[Images](#)

[News](#)

[Maps](#)

[More](#)

[Settings](#)

[Tools](#)

About 57,900 results (0.57 seconds)

Filtering Data with dplyr - learn data science

<https://blog.exploratory.io/filter-data-with-dplyr-76cf5f1a258e>

Of course, dplyr has 'filter()' function to do such filtering, but there is even more. ... get used to it especially if you're coming from outside of R world, but you are going ... closer you'd notice that there are some NA values in ARR_DELAY column.

Removing NA in dplyr pipe - Stack Overflow

<https://stackoverflow.com/questions/removing-na-in-dplyr-pipe>

1 answer

Nov 1, 2014 - If you just want to remove NA s, use na.omit (base) or tidyrr::drop_na : ... If you only want to remove NA s from the HeartAttackDeath column, filter with is.na , or use ... but it's a bit trickier to put in a chain because it takes a data frame as an ... From a fresh R session
library(dplyr); library(hflights); x = hflights ...

Removing NA observations with dplyr::filter()

2 answers Mar 4, 2015

dplyr: Filter multiple conditions with **selection NA ...

2 answers Aug 25, 2017

Filter based on NA in dplyr

1 answer Jan 16, 2015

Remove rows where all variables are NA using dplyr

3 answers Jan 18, 2017

[More results from stackoverflow.com](#)

Removing NA observations with dplyr::filter() - Stack Overflow

<https://stackoverflow.com/questions/removing-na-observations-with-dp...>

2 answers

Mar 4, 2015 - R doesn't know what you are doing in your analysis, so instead of potentially introducing bugs that would later end up being published an ...

Getting help



[r] tidyverse how do i delete column



[All](#) [Videos](#) [Shopping](#) [News](#) [Images](#) [More](#) [Settings](#) [Tools](#)

About 50,900 results (0.57 seconds)

How to drop column in R?

Select **columns** by typing their names

You can use "-" (minus) to **drop columns**. All you need to do is to add '-' (minus) right before the **columns** you want to **drop**. It's that simple. Notice that the last **column** name inside the 'select()' function where I'm using "" (back-tick) to surround "NA" characters. Mar 10, 2016

[Selecting columns and renaming are so easy with dplyr](#)

<https://blog.exploratory.io/selecting-columns-809bdd1ef615>

Search for: [How to drop column in R?](#)

[About](#) [Featured Snippets](#) [Feedback](#)

People also ask

How do I remove a column from a dataset in R? ▾

How do I drop a column in R? ▾

How do I rename a column in Tidyverse? ▾

How do I select a column in Dplyr? ▾

[Feedback](#)

Drop column in R using Dplyr - drop variables - DataScience ...

www.datasciencemadesimple.com/drop-variables-columns-r-using-dplyr ▾

Drop column in R using Dplyr: Drop column in R can be done by using minus before the select function. Dplyr package in R is provided with select() function which is used to select or drop the columns based on conditions.

Getting help



[r] tidyverse how do i fit regression



All

Videos

Images

Shopping

News

More

Settings

Tools

About 33,700 results (0.57 seconds)

Tidyverse-Friendly Introductory Linear Regression • moderndive

<https://moderndive.github.io> › [moderndive](#) ▼

An R package of datasets and wrapper functions for tidyverse-friendly introductory linear regression used in. ModernDive: An Introduction to Statistical and Data ...

Linear regression | Computing for the Social Sciences

<https://cfss.uchicago.edu> › [notes](#) › [linear-models](#) ▼

library(tidyverse) library(modelr) library(broom) library(rcfss) set.seed(1234) ... R for Data Science walks you through the steps to perform all these calculations ...

broom and dplyr - CRAN

<https://cran.r-project.org> › [web](#) › [packages](#) › [broom](#) › [vignettes](#) › [broom_a...](#) ▼

Apr 7, 2019 - Often, we want to perform multiple tests or fit multiple models, each on a ... This workflow becomes even more useful when applied to regressions. ... R-squared: 0.8345, Adjusted R-squared: 0.8295 ## F-statistic: 166.4 on 1 ...

Introduction to broom - CRAN

<https://cran.r-project.org> › [vignettes](#) › [broom](#) ▼

Apr 7, 2019 - The broom package takes the messy output of built-in functions in R, such as ... This includes coefficients and p-values for each term in a regression, ... Instead, you can use the tidy function, from the broom package, on the fit.

Filtering data

We already know one way to filter data:

```
is_tall <- fev_data$height > 72  
fev_data[is_tall, ]
```

```
## # A tibble: 7 x 7  
##   seqnbr subjid   age   fev height  sex smoke  
##   <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1    401  18841    14  4.27  72.5     1     0  
## 2    450  32741    13  4.22   74     1     0  
## 3    464  37241    13  4.88   73     1     0  
## 4    517  49541    13  5.08   74     1     0  
## 5    550  59941    14  4.27  72.5     1     0  
## 6    632  37441    17  5.63   73     1     0  
## 7    636  44241    16  3.64  73.5     1     0
```

Let's learn another way!

Filtering data

Why learn another way?

Easier to remember, faster to write, more robust to errors, highly scalable. . .

If you don't believe me by the end of the hour, let's chat about it after class!

Filtering data

```
fev_data %>%  
  filter(height > 72)
```

```
## # A tibble: 7 x 7
```

##	seqnbr	subjid	age	fev	height	sex	smoke
##	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	401	18841	14	4.27	72.5	1	0
## 2	450	32741	13	4.22	74	1	0
## 3	464	37241	13	4.88	73	1	0
## 4	517	49541	13	5.08	74	1	0
## 5	550	59941	14	4.27	72.5	1	0
## 6	632	37441	17	5.63	73	1	0
## 7	636	44241	16	3.64	73.5	1	0

Filtering data

```
fev_data %>%  
  filter(age == 6)
```

```
## # A tibble: 37 x 7
```

```
##   seqnbr subjid   age   fev height  sex smoke  
##   <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1      7   1752     6  1.92    58     0     0  
## 2      8   1753     6  1.42    56     0     0  
## 3     11   1952     6  1.60    53     0     0  
## 4     18   3551     6  1.88    53     0     0  
## 5     49  10841     6  1.65    55     1     0  
## 6     55  12241     6  1.63    54     1     0  
## 7     63  14251     6  1.48    51     0     0  
## 8     66  14541     6  1.75   57.5     1     0  
## 9     80  16151     6  1.72    53     0     0  
## 10    82  16252     6  1.70    53     0     0  
## # ... with 27 more rows
```

Filtering data

```
fev_data %>%  
  filter(age != 20)
```

```
## # A tibble: 654 x 7
```

```
##   seqnbr subjid   age   fev height  sex smoke  
##   <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1      1    301     9  1.71   57     0     0  
## 2      2    451     8  1.72  67.5     0     0  
## 3      3    501     7  1.72  54.5     0     0  
## 4      4    642     9  1.56   53     1     0  
## 5      5    901     9  1.90   57     1     0  
## 6      6   1701     8  2.34   61     0     0  
## 7      7   1752     6  1.92   58     0     0  
## 8      8   1753     6  1.42   56     0     0  
## 9      9   1901     8  1.99  58.5     0     0  
## 10     10   1951     9  1.94   60     0     0  
## # ... with 644 more rows
```

Filtering data

Syntax for filtering rules for continuous data:

- ▶ `age == 6`
- ▶ `age != 6`
- ▶ `fev > 2`
- ▶ `fev <= 1.5`: less than or equal to
- ▶ `is.na(age)`
- ▶ `!is.na(age)`

Filtering data

Syntax for filtering rules for categorical data:

- ▶ `sex == "F" or sex != "F"`
- ▶ `sex %in% c("M", "F")`

Filtering data

You can filter based on multiple criteria, too!

```
fev_data %>%  
  filter(age == 14, smoke != 0)
```

```
## # A tibble: 7 x 7  
##   seqnbr subjid   age   fev height  sex smoke  
##   <dbl>  <dbl> <dbl> <dbl>  <dbl> <dbl> <dbl>  
## 1    332   4952    14  2.24    66     0     1  
## 2    358  10053    14  3.43    64     0     1  
## 3    370  11642    14  3.96    72     1     1  
## 4    384  15751    14  3.07    65     0     1  
## 5    439  30042    14  4.31    69     1     1  
## 6    556  61941    14  2.28    66     1     1  
## 7    602  82743    14  4.76    68     1     1
```


Selecting columns

Now we know how to select rows. How do we select columns?

```
fev_data %>%  
  select(fev, height, age)
```

```
## # A tibble: 654 x 3  
##       fev height  age  
##   <dbl> <dbl> <dbl>  
## 1  1.71    57     9  
## 2  1.72   67.5    8  
## 3  1.72   54.5    7  
## 4  1.56    53     9  
## 5  1.90    57     9  
## 6  2.34    61     8  
## 7  1.92    58     6  
## 8  1.42    56     6  
## 9  1.99   58.5    8  
## 10 1.94    60     9  
## # ... with 644 more rows
```

Selecting columns

We can also drop columns

```
fev_data %>%  
  select(-seqnbr, -subjid)
```

```
## # A tibble: 654 x 5  
##       age    fev height    sex smoke  
##   <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1     9  1.71    57      0     0  
## 2     8  1.72   67.5     0     0  
## 3     7  1.72   54.5     0     0  
## 4     9  1.56    53     1     0  
## 5     9  1.90    57     1     0  
## 6     8  2.34    61     0     0  
## 7     6  1.92    58     0     0  
## 8     6  1.42    56     0     0  
## 9     8  1.99   58.5     0     0  
## 10    9  1.94    60     0     0  
## # ... with 644 more rows
```

Summarising data

Now that we know how to subset our data, how do we summarize it?

```
fev_data %>%  
  filter(age == 14, smoke != 0) %>%  
  summarize(mean(fev))
```

```
## # A tibble: 1 x 1  
##   `mean(fev)`  
##         <dbl>  
## 1         3.43
```

Summarising data

Now that we know how to subset our data, how do we summarize it?

```
fev_data %>%  
  filter(age == 14, smoke != 0) %>%  
  summarize(my_mean = mean(fev))
```

```
## # A tibble: 1 x 1  
##   my_mean  
##   <dbl>  
## 1     3.43
```

You can give a summary variable a name, e.g., `my_mean`

Summarising data

We can summarize in multiple ways at once!

```
fev_data %>%  
  filter(age == 14, smoke != 0) %>%  
  summarize(mean(fev), sd(fev))
```

```
## # A tibble: 1 x 2  
##   `mean(fev)` `sd(fev)`  
##           <dbl>      <dbl>  
## 1         3.43      0.976
```

Both summarise and summarize work!

Summarising data

But what if I want the average FEV for both smokers and non-smokers? Do I have to repeat this?

No – I can create a grouping variable!

Summarising data

Create a grouping variable with `group_by`

```
fev_data %>%  
  group_by(smoke)
```

```
## # A tibble: 654 x 7  
## # Groups:   smoke [2]  
##   seqnbr subjid   age   fev height   sex smoke  
##   <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1      1      301     9  1.71    57      0     0  
## 2      2      451     8  1.72   67.5      0     0  
## 3      3      501     7  1.72   54.5      0     0  
## 4      4      642     9  1.56    53      1     0  
## 5      5      901     9  1.90    57      1     0  
## 6      6     1701     8  2.34    61      0     0  
## 7      7     1752     6  1.92    58      0     0  
## 8      8     1753     6  1.42    56      0     0  
## 9      9     1901     8  1.99   58.5      0     0  
## 10     10     1951     9  1.94    60      0     0  
## # ... with 644 more rows
```

This doesn't change the dataset in any way (except the groups are listed)

Summarising data

Now when we go to summarize there is a summary for each group:

```
fev_data %>%  
  group_by(smoke) %>%  
  summarize(mean(fev), sd(fev))
```

```
## # A tibble: 2 x 3  
##   smoke `mean(fev)` `sd(fev)`  
##   <dbl>      <dbl>      <dbl>  
## 1     0         2.57         0.851  
## 2     1         3.28         0.750
```

Note: smoke == 1 is smokers. Interesting...

Summarising data

It's hard to interpret differences in mean and standard deviation without knowledge of the sample size in each group!

```
fev_data %>%  
  group_by(smoke) %>%  
  summarize(mean(fev), sd(fev))
```

```
## # A tibble: 2 x 3  
##   smoke `mean(fev)` `sd(fev)`  
##   <dbl>         <dbl>         <dbl>  
## 1      0          2.57          0.851  
## 2      1          3.28          0.750
```

Summarising data

Use `n()` with `summarize` to get the number of observations in each group:

```
fev_data %>%  
  group_by(smoke) %>%  
  summarize(n = n(), mean = mean(fev), sd = sd(fev))
```

```
## # A tibble: 2 x 4  
##   smoke      n mean    sd  
##   <dbl> <int> <dbl> <dbl>  
## 1     0   589  2.57 0.851  
## 2     1    65  3.28 0.750
```

Summarising data

You can also summarize based on your own variables!

```
fev_data %>%  
  group_by(height < 60) %>%  
  summarize(n(), mean(fev))
```

```
## # A tibble: 2 x 3  
##   `height < 60` `n()` `mean(fev)`  
##   <lgl>         <int>      <dbl>  
## 1 FALSE         409        3.10  
## 2 TRUE          245        1.86
```

Think-pair-share

```
fev_data$sex %>% unique
```

```
## [1] 0 1
```

Suppose I was given this dataset without information about variable coding. We have 2 sexes observed in this dataset – most likely binary male and female. How could we figure out if `sex == 1` is male or female?

```
fev_data %>%  
  select(-seqnbr, -subjid) %>%  
  names
```

```
## [1] "age"      "fev"      "height"   "sex"      "smoke"
```

Discuss with a partner how you might investigate this (2 minutes).

Think-pair-share

[Live demo]

Final data analysis

Solving the smokers-have-better-lung-function paradox:

```
fev_data %>%  
  group_by(smoke) %>%  
  summarize(n(), mean(fev), sd(fev))
```

```
## # A tibble: 2 x 4  
##   smoke `n()` `mean(fev)` `sd(fev)`  
##   <dbl> <int>      <dbl>      <dbl>  
## 1     0   589        2.57        0.851  
## 2     1    65        3.28        0.750
```

This was surprising to me. What could be confounded with smoking status?

Final data analysis: smoking and lung function in children

```
fev_data %>%  
  group_by(smoke) %>%  
  summarize(mean(age))
```

```
## # A tibble: 2 x 2  
##   smoke `mean(age)`  
##   <dbl>      <dbl>  
## 1     0        9.53  
## 2     1       13.5
```

Smokers are typically *older* than non-smokers!

Final data analysis: smoking and lung function in children

```
fev_data %>%  
  group_by(age, smoke) %>%  
  summarize(n(), mean(fev), sd(fev))
```

```
## # A tibble: 28 x 5  
## # Groups:   age [17]  
##       age smoke `n()` `mean(fev)` `sd(fev)`  
##   <dbl> <dbl> <int>      <dbl>      <dbl>  
## 1     3     0     2      1.24      0.235  
## 2     4     0     9      1.28      0.353  
## 3     5     0    28      1.55      0.308  
## 4     6     0    37      1.66      0.223  
## 5     7     0    54      1.87      0.335  
## 6     8     0    85      2.12      0.391  
## 7     9     0    93      2.44      0.478  
## 8     9     1     1      1.95      NA  
## 9    10     0    76      2.67      0.533  
## 10    10     1     5      3.06      0.441  
## # ... with 18 more rows
```


Sorting data

```
fev_data %>%  
  group_by(age, smoke) %>%  
  summarize(n(), mean(fev), sd(fev)) %>%  
  arrange(age) # arrange by increasing age
```

```
## # A tibble: 28 x 5  
## # Groups:   age [17]  
##   age smoke `n()` `mean(fev)` `sd(fev)`  
##   <dbl> <dbl> <int>      <dbl>      <dbl>  
## 1     3     0     2       1.24       0.235  
## 2     4     0     9       1.28       0.353  
## 3     5     0    28       1.55       0.308  
## 4     6     0    37       1.66       0.223  
## 5     7     0    54       1.87       0.335  
## 6     8     0    85       2.12       0.391  
## 7     9     0   93       2.44       0.478  
## 8     9     1     1       1.95      NA  
## 9    10     0    76       2.67       0.533  
## 10   10     1     5       3.06       0.441  
## # ... with 18 more rows
```

Putting it all together

```
fev_data %>%  
  group_by(age, smoke) %>%  
  summarize(n = n(), # name summary statistics columns...  
            mean = mean(fev), sd = sd(fev)) %>%  
  filter(n >= 5) %>% # ... so we can filter on them later  
  arrange(desc(age)) # arrange by decreasing age
```

```
## # A tibble: 21 x 5  
## # Groups:   age [14]  
##       age smoke      n  mean    sd  
##   <dbl> <dbl> <int> <dbl> <dbl>  
## 1     17      0      6  4.65 0.871  
## 2     16      0      6  3.75 0.567  
## 3     16      1      7  3.60 0.858  
## 4     15      0      9  3.92 0.960  
## 5     15      1     10  3.09 0.766  
## 6     14      0     18  3.64 0.627  
## 7     14      1      7  3.43 0.976  
## 8     13      0     30  3.53 0.717  
## 9     13      1     13  3.38 0.776  
## 10     12      0     50  3.24 0.733  
## # ... with 11 more rows
```

Summary

Today we learnt 5 tidyverse functions that facilitate data analysis:

- ▶ `filter()` picks rows based on their values
- ▶ `select()` picks columns based on their names
- ▶ `summarize()` reduces multiple values to a single summary
- ▶ `group_by()` groups rows together (useful for `summarize`)
- ▶ `arrange()` changes the ordering of the rows

The plan

- ▶ 5 minute break
- ▶ In-class exercise available via Canvas and github.com/adw96/biost509
 - ▶ Designed to be completed by 3:20 p.m.
 - ▶ Due today 5 p.m.
 - ▶ Place a *yellow sticky note* on your computer to indicate you are stuck, or a *blue sticky note* to indicate you have a non-urgent question
- ▶ Homework due next week by 1 p.m. Friday

The TAs and I will walk around to help you; we will help *yellow stickies*, then *blue stickies*. Please put your sticky note high enough that we can see it!